

# IEMS 5709 Spring 2016 Homework 2

Release date: Mar 2, 2016

Due date: 23:59, Mar 21, 2016

*The solution will be posted right after the deadline, so no late homework will be accepted!*

Every Student **MUST** include the following statement, together with his/her signature in the submitted homework.

*I declare that the assignment submitted on Elearning system is original except for source material explicitly acknowledged, and that the same or related material has not been previously submitted for another course. I also acknowledge that I am aware of University policy and regulations on honesty in academic work, and of the disciplinary guidelines and procedures applicable to breaches of such policy and regulations, as contained in the website*

<http://www.cuhk.edu.hk/policy/academichonesty/>.

Signed (Student \_\_\_\_\_) Date: \_\_\_\_\_

Name \_\_\_\_\_ SID \_\_\_\_\_

## Submission notice:

- Submit your homework via the elearning system

## General homework policies:

A student may discuss the problems with others. However, the work a student turns in must be created COMPLETELY by oneself ALONE. A student may not share ANY written work or pictures, nor may one copy answers from any source other than one's own brain.

Each student **MUST LIST** on the homework paper the **name of every person he/she has discussed or worked with**. If the answer includes content from any other source, the student **MUST STATE THE SOURCE**. Failure to do so is cheating and will result in sanctions. Copying answers from someone else is cheating even if one lists their name(s) on the homework.

If there is information you need to solve a problem but the information is not stated in the problem, try to find the data somewhere. If you cannot find it, state what data you need, make a reasonable estimate of its value, and justify any assumptions you make. You will be graded not only on whether your answer is correct, but also on whether you have done an intelligent analysis.

## Q1 [40 marks]: PageRank Algorithm on GraphLab and Hadoop

In this question, you are required to run the PageRank algorithm on the given dataset, and then output the top 100 nodes ranked by their PageRank scores. The dataset is a directed web graph from the Google programming contest, consisting of 875,713 nodes and 5,105,039 edges. See Ref [1] for more details about the data statistics. Note that the data has been pre-processed to exclude all the dead-end nodes and it is available in Ref [2]. Each line of the data is a [source ID] [target ID] pair separated by whitespace.

**(a) [10 marks]** GraphLab Create [1] is an extensible machine learning framework that enables developers and data scientists to easily build and deploy intelligent applications and services at scale. Follow the procedure in Ref [3] to install GraphLab Create on one single virtual machine. In your homework submission, show the commands you used.

**(b) [10 marks]** On one machine, run the PageRank toolkit, provided by GraphLab Create in Ref [4]. Submit the code and the top 100 nodes.

**(c) [20 marks]** Dato Distributed (Ref [5]) provides a lightweight framework for creating an environment for distributed execution and submit jobs to these environments. Install Dato Distributed on 4 virtual machines and run the PageRank algorithm in a Dato Distributed cluster. Ref [6] shows how to setup a Dato Distributed cluster and Ref [7] shows how to run PageRank in a Dato Distributed cluster.

Please change the threshold for convergence and maximum iterations in the PageRank toolkit and run at least three times with different parameters. Compare and explain the results. Submit the code and compare your results of your different runs by listing the corresponding top 100 nodes, number of iterations and time-to-convergence/completion.

**Hints:** The PageRank toolkit provides default values for those parameters. You can tune around those default values.

## Q2 [60 marks]: Word Counting on a Storm Cluster

In this question, you are required to implement the word-counting algorithm under the Storm platform. The dataset is available at:

<https://github.com/YangRonghai/100ebooks/blob/master/StormData?raw=true>

The basic idea is to first define a `FileReaderSpout` class to read an input file and emit a sentence every time. Then you can implement a `SplitSentenceBolt` class to split every sentence into words and emit the word to the next `WordCountBolt` class, which needs to compute the frequency of each word. The word count result should be dumped to a persistent storage. You are allowed to refer to (or even borrow codes) from publicly available Storm examples/source-codes AS LONG AS you clearly list and acknowledge your sources in your submission.

### (a) [15 marks] Apache Storm Cluster Setup.

Setup an Apache Storm cluster with 4 VMs using Amazon EC2, or Google Compute Engine, or Windows Azure or the IE DIC. The deployment of a Storm cluster is shown in Fig. 1. Basically, there are three workers and one Nimbus worker. The Nimbus worker also runs the ZooKeeper service and Storm UI service. Ref [8] and [9] show how to install a distributed Apache Storm cluster. You may need to modify some configurations to suite your own computing environment.

You are suggested to reuse your cluster in the Homework#1 to save your time.

There are some examples under the directory of `Storm_HOME/examples/storm-starter`. You are required to run the example `ExclamationTopology` to prove your installation to be successful. You need to show the key steps of the installation as well as the running results, including the snapshot of the Storm UI and `Storm_HOME/logs/`.

### Hints:

- You ARE NOT allowed to use the [storm-deploy](#) project to provision Storm clusters on AWS, namely, you need to manually configure the Storm system.
- Use the latest version storm-0.10.0  
<http://apache.01link.hk/storm/apache-storm-0.10.0/apache-storm-0.10.0.tar.gz>
- To facilitate debugging, you can find the logs under the `Storm_HOME/logs/` directory upon any errors.

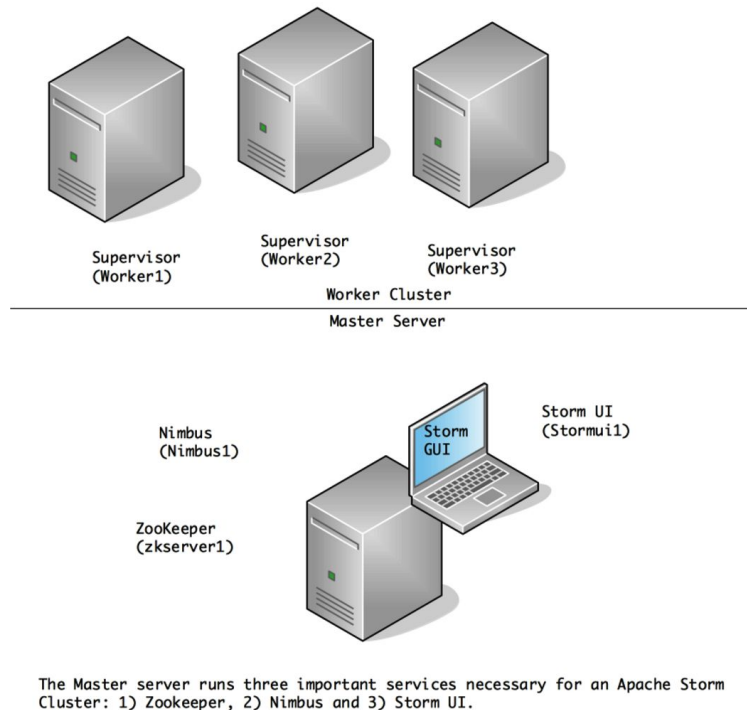


Fig. 1 Cluster Topology

### (b) [15 marks] Find Frequent Words

Write a word counting program without any processing guarantee, i.e. under the at-most-once model. Your program should read the input file, count the frequency of words and dump the final result to the persistent storage (e.g., Linux file system or HDFS). Storm utilizes Maven [10] to compile and build the executable Jar. You can reuse the same configuration file (i.e., pom.xml) under the *Storm\_HOME/examples/storm-starter/* directory. Find out the Top 10 most frequently used words in the dataset and their corresponding frequency. Submit your code and the final top 10 words and their corresponding frequency.

#### Hints:

- You are allowed to post-process the result dumped from Storm.
- DO NOT emit the entire file in one `nextTuple()` to avoid failure caused by network congestion. You can emit one line for each `nextTuple()` call.
- In case of buffer overflow, you can configure `topology.max.spout.pending` [14] to specify the maximum number of unacknowledged tuples that can exist in your topology (from that spout) at a given time. Otherwise, there would be lots of failures given the relatively limited resource in your small cluster.
- Note that when you read a file in the open method, each worker would go to its local path to find the file. With this method, you need to copy the input file to every worker node. A better (yet more complicated) option to solve this issue is to use the shared HDFS. Both ways are acceptable.
- Copy your java code to the following directory:  
*Storm\_HOME/examples/storm-starter/src/jvm/storm/starter/*  
 and then use Maven to compile.

**(c) [15 marks] Guaranteed Message/Tuple Processing: the At-Least-Once Model.**

Refer to [11,12,13] to change the program in part (b), so that the topology can guarantee to process every message (i.e., at least once model). You are required to manually fail the delivery of “the” for 10 times. In other words, you can use *collector.fail(tuple)* to simulate the error in the first bolt when the word “the” appears for the first 10 times. Under the at-least-once model, find out the Top 10 most frequently used words. With these artificial errors, re-run the program in part (b) and output the Top 10 most frequently used words. Do they generate the same result? Why or why not? Compare the performance of these two models (including running time, number of tuples emitted, number of tuples Acked). Submit the code and the results. Explain your findings.

**Hints:**

- You can use the Java built-in function, e.g., *System.currentTimeMillis()*, to get the running time.

**(d) [15 marks] Scaling the topology.**

Refer to Page 41 of the lecture slides on Storm to change the parallelism degree (i.e., # of Spout threads, # of Bolt threads) of your program. Re-run the program in part (b) and compare the performance (e.g., running time, number of tuples emitted, number of tuples Acked, number of tuples Failed) under the following 4 different settings. Explain your findings.

	# of FileReaderSpout	# of SplitSentenceBolt	# of WordCountBolt
Setting 1	3	3	3
Setting 2	1	3	3
Setting 3	1	1	3
Setting 4	3	10	10

## Reference:

- [1] SNAP Google Web Data  
<https://snap.stanford.edu/data/web-Google.html>
- [2] Processed Google Web Graph Data  
[https://googledrive.com/host/0B1\\_aqhmFB5SFQTdNRURBWXYZyZzQ/web-Google.txt](https://googledrive.com/host/0B1_aqhmFB5SFQTdNRURBWXYZyZzQ/web-Google.txt)
- [3] Installation Guide for GraphLab Create  
<https://dato.com/download/install-graphlab-create.html>
- [4] GraphLab Create PageRank Toolkit  
[https://dato.com/products/create/docs/graphlab.toolkits.graph\\_analytics.html#pagerank](https://dato.com/products/create/docs/graphlab.toolkits.graph_analytics.html#pagerank)
- [5] Dato Distributed Introduction  
<https://dato.com/learn/userguide/deployment/pipeline-introduction.html>
- [6] Dato Distributed Setup  
<https://dato.com/learn/userguide/deployment/pipeline-hadoop-setup.html>
- [7] Distributed Machine Learning in Dato Distributed Cluster (Follow the Hadoop section)  
<https://dato.com/learn/userguide/deployment/pipeline-dml.html>
- [8] How to install a distributed apache Storm cluster  
<http://knowm.org/how-to-install-a-distributed-apache-storm-cluster/>
- [9] Setting up a Storm Cluster  
<http://storm.apache.org/documentation/Setting-up-a-Storm-cluster.html>
- [10] Storm Maven  
<http://storm.apache.org/documentation/Maven.html>
- [11] Guaranteeing message processing:  
<https://storm.apache.org/documentation/Guaranteeing-message-processing.html>
- [12] At least once example:  
<https://www.safaribooksonline.com/library/view/getting-started-with/9781449324025/ch04.html>
- [13] Storm fault tolerance in practice  
<https://www.anchormen.nl/apache-storms-fault-tolerance-in-practice/>
- [14] Storm configuration  
<http://storm.apache.org/documentation/Configuration.html>